Data Summarization

Introduction to R for Public Health Researchers

Data Summarization

- Basic statistical summarization
 - mean (x): takes the mean of x
 - sd(x): takes the standard deviation of x
 - median (x): takes the median of x
 - quantile(x): displays sample quantities of x. Default is min, IQR, max
 - range (x): displays the range. Same as c(min(x), max(x))
 - sum(x):Sum of X
- Transformations
 - log log (base e) transformation
 - log2 log base 2 transform
 - log10 log base 10 transform
 - sqrt square root

Some examples

We can use the mtcars to explore different ways of summarizing data. The head command displays the first 6 (default) rows of an object:

head (mtcars)

	mpg	cyl	disp	hp	drat	wt	qsec	VS	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46		1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02		1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1		3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02			3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1		3	1

Note - the \$ references/selects columns from a data.frame/tibble:

mean (mtcars\$hp)

[1] 146.6875

quantile (mtcars\$hp)

 0%
 25%
 50%
 75%
 100%

 52.0
 96.5
 123.0
 180.0
 335.0

```
median(mtcars$wt)

[1] 3.325

quantile(mtcars$wt, probs = 0.6)

60%
3.44
```

t.test will be covered more in detail later, gives a mean and 95% CI:

```
t.test(mtcars$wt)

One Sample t-test

data: mtcars$wt
t = 18.6, df = 31, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
   2.864478 3.570022
sample estimates:
mean of x
   3.21725</pre>
```

Note that many of these functions have additional inputs regarding missing data, typically requiring the na.rm argument ("remove NAs").

```
x = c(1,5,7,NA,4,2,8,10,45,42)
mean(x)

[1] NA

mean(x, na.rm = TRUE)

[1] 13.77778

quantile(x, na.rm = TRUE)

0% 25% 50% 75% 100%
1 4 7 10 45
```

Data Summarization on matrices/data frames

- Basic statistical summarization
 - rowMeans (x): takes the means of each row of x
 - colMeans (x): takes the means of each column of x
 - rowSums(x):takes the sum of each row of x
 - colSums (x): takes the sum of each column of x
 - summary (x): for data frames, displays the quantile information

Lab Part 1

Website

TB Incidence

Please download the TB incidence data:

http://johnmuschelli.com/intro_to_r/data/tb_incidence.xlsx

Here we will read in a data.frame of values from TB incidence:

```
# A tibble: 6 x 19
  TB incidence, all forms (per 100 000 population per year)
  ... with 17 more variables: `1991` <dbl>, `1992` <dbl>, `1993` <dbl>,
    `1994` <dbl>, `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>,
   `1999` <dbl>, `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>,
          <dbl>, `2005` <dbl>, `2006` <dbl>, `2007` <dbl>
```

Indicator of TB

We can rename the first column to be the country measured using the rename function in dplyr (we have to use the `things because there are spaces in the name):

colnames will show us the column names and sho that country is renamed:

```
[1] "country" "1990"
                          "1991"
                                    "1992"
                                               "1993"
                                                         "1994"
                                                                    "1995"
[8] "1996"
               "1997"
                                    "1999"
                                               "2000"
                                                         "2001"
                                                                    "2002"
[15] "2003"
              "2004"
                          "2005"
                                    "2006"
                                               "2007"
```

Column and Row means

colMeans and rowMeans must work on all numeric data. We will subset years before 2000 (starting with 1):

```
105.5797 107.6715 108.3140 110.3188 111.9662 114.1981 115.3527 118.8792
121.5169 125.0435
head(tb[, c("country", "before 2000 avg")])
                           168.0
                           41.8
                           8.5
```

Summary

Using summary can give you rough snapshots of each column, but you would likely use mean, min, max, and quantile when necessary:

summary (tb)

```
Min. : 0.0
                            Min. : 4.0
                                         Min. : 2.0
Class : character
               1st Qu.: 27.5
                                         1st Qu.: 27.0
Mode : character
               Median: 60.0
                                         Median: 56.0
                            Mean :107.7
               Mean :105.6
                                         Mean :108.3
                                         3rd Qu.:171.5
               3rd Qu.:165.0
               Max. :585.0
                                         Max. :606.0
                            NA's :1
                                         NA's :1
               NA's :1
            Min. : 0 Min. : 3.0 Min. : 0.0
Min. : 4.0
1st Qu.: 27.5
                                     1st Qu.: 25.5
            1st Qu.: 26 1st Qu.: 26.5
            Median: 57 Median: 58.0
                                     Median: 60.0
Median: 56.0
Mean :110.3
                                     Mean :115.4
3rd Qu.:171.0
                       3rd Qu.:177.5
                                      3rd Qu.:179.0
Max. :618.0
            Max. :630 Max. :642.0
                                     Max. :655.0
             NA's :1
NA's :1
                                     NA's
                        NA's
Min. : 0.0
             Min. : 0.0 Min. : 0.0 Min. : 0.0
            1st Qu.: 23.5 1st Qu.: 22.5 1st Qu.: 21.5
```

Apply statements

You can apply more general functions to the rows or columns of a matrix or data frame, beyond the mean and sum.

```
apply(X, MARGIN, FUN, ...)
```

X : an array, including a matrix.

MARGIN: a vector giving the subscripts which the function will be applied over. E.g., for a matrix 1 indicates rows, 2 indicates columns, c(1, 2) indicates rows and columns. Where X has named dimnames, it can be a character vector selecting dimension names.

FUN: the function to be applied: see 'Details'.

...: optional arguments to FUN.

Apply statements

```
105.5797 107.6715 108.3140 110.3188 111.9662 114.1981 115.3527 118.8792
121.5169 125.0435
110.6440 112.7687 114.4853 116.6744 120.0931 122.7119 126.1800 131.0858
137.3754 146.0755
```

Other Apply Statements

- tapply(): 'grouping' apply
- lapply(): 'list' apply [tomorrow]
- sapply(): 'simple' apply [tomorrow]
- Other less used ones...

See more details here: http://nsaunders.wordpress.com/2010/08/20/a-brief-introduction-to-apply-in-r/

Youth Tobacco Survey

Please download the Youth Tobacco Survey data. You can also read it in directly from the web:

```
# A tibble: 6 x 31
 ... with 27 more variables: TopicDesc <chr>, MeasureDesc <chr>,
   DataSource <chr>, Response <chr>, Data Value Unit <chr>,
   Data Value Type <chr>, Data Value <dbl>,
   Data Value Footnote Symbol <chr>, Data Value Footnote <chr>,
   Data Value Std Err <dbl>, Low Confidence Limit <dbl>,
   High Confidence Limit <dbl>, Sample Size <int>, Gender <chr>,
   Race <chr>, Age <chr>, Education <chr>, GeoLocation <chr>,
   TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>,
```

Length and unique

unique (x) will return the unique elements of x

unique (yts\$LocationDesc) [1:10]

```
[1] "Arizona" "Connecticut"
[3] "Georgia" "Hawaii"
[5] "Illinois" "Louisiana"
[7] "Mississippi" "Utah"
[9] "Missouri" "National (States and DC)"
```

length will tell you the length of a vector. Combined with unique, tells you the number of unique elements:

```
length(unique(yts$LocationDesc))
```

[1] 50

Table

table (x) will return a frequency table of unique elements of x

table (yts\$LocationDesc) [1:5]

Alabama Arizona Arkansas California Colorado 378 240 210 96 48

Lab Part 2

Website

Subsetting to specific columns

Let's just take smoking status measures for all genders using filter, and the columns that represent the year, state using select:

Perform Operations By Groups: dplyr

group_by allows you group the data in a more intuitive way than tapply

We will use group_by to group the data by line, then use summarize (or summarise) to get the mean percentage of current smokers:

```
summarize(group by(sub yts, YEAR), year avg = mean(Data Value, na.rm = TRUE))
# A tibble: 17 x 2
 2 2000 19.878431
 3 2001 15.661111
 6 2004 13 926923
   2007 13.013636
   2009 11.663333
   2010 12.290000
   2011 11.773913
   2012 9.954545
```

Using the pipe (comes with dplyr):

Pipe sub_yts into group_by, then pipe that into summarize:

Counting

Standard statistics can be calculated. There are other functions, such as n () count the number of observations, tally () to count as a wrapper:

```
# A tibble: 6 x 2
```

Lab Part 3

Website

Data Summarization/Visualization: ggplot2

ggplot2 is a package of plotting that is very popular and powerful (using the grammar of graphics). We will use qplot ("quick plot") for most of the basic examples:

aplot

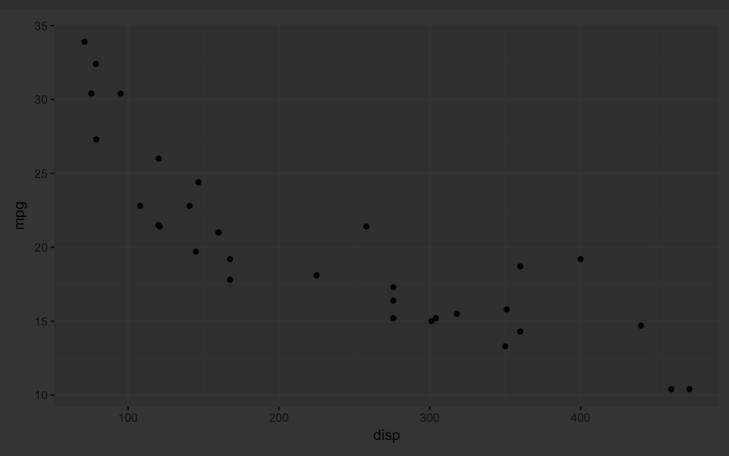
```
function (x, y = NULL, ..., data, facets = NULL, margins = FALSE,
   geom = "auto", xlim = c(NA, NA), ylim = c(NA, NA), log = "",
   main = NULL, xlab = deparse(substitute(x)), ylab = deparse(substitute(y)),
   asp = NA, stat = NULL, position = NULL)
       warning("`stat` is deprecated", call. = FALSE)
       warning("`position` is deprecated", call. = FALSE)
       stop("'geom' must be a character vector", call. = FALSE)
```

Basic Plots

Plotting is an important component of exploratory data analysis. We will review some of the more useful and informative plots here. We will go over formatting and making plots look nicer in additional lectures.

Scatterplot

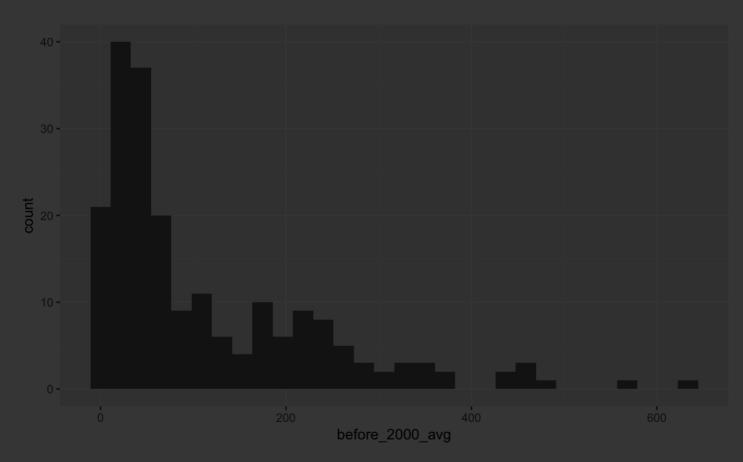
```
library(ggplot2)
qplot(x = disp, y = mpg, data = mtcars)
```



Histograms

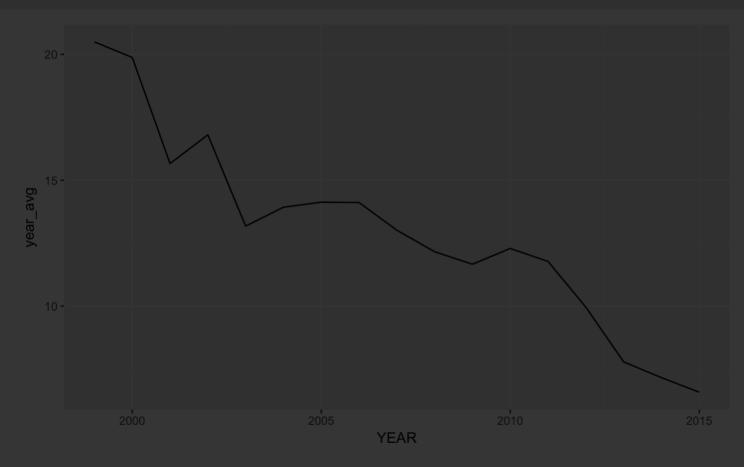
qplot(x = before_2000_avg, data = tb, geom = "histogram")

Warning: Removed 1 rows containing non-finite values (stat_bin).



Plot with a line

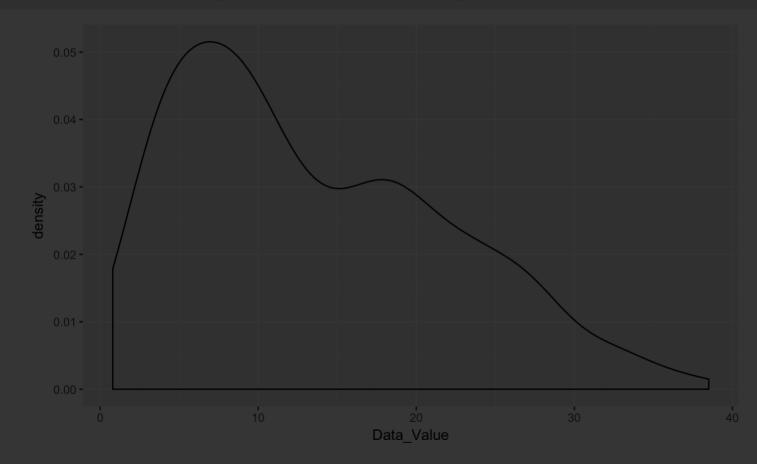
qplot(x = YEAR, y = year_avg, data = yts_avgs, geom = "line")



Density

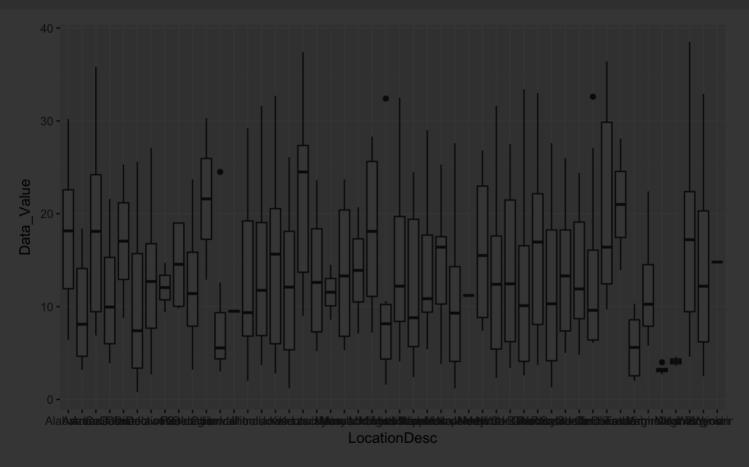
Over all years and states, this is the density of smoking status incidence:

qplot(x = Data_Value, data = sub_yts, geom = "density")

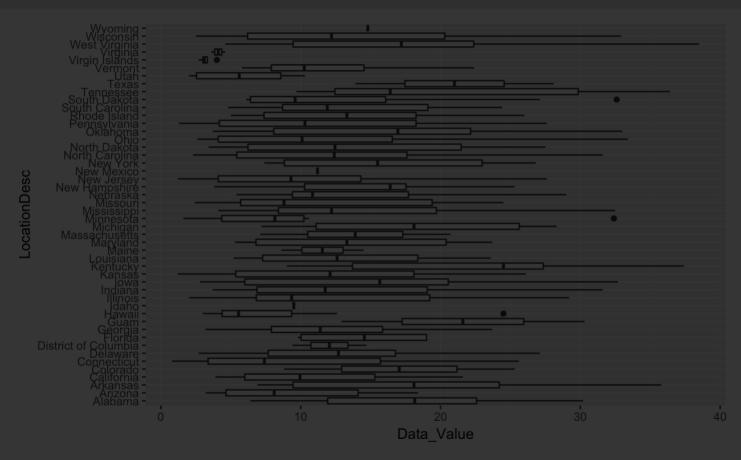


Boxplots

qplot(x = LocationDesc, y = Data_Value, data = sub_yts, geom = "boxplot")



Boxplots



Data Summarization for data.frames

- Basic summarization plots
 - matplot(x,y): scatterplot of two matrices, x and y
 - pairs (x, y): plots pairwise scatter plots of matrices x and y, column by column

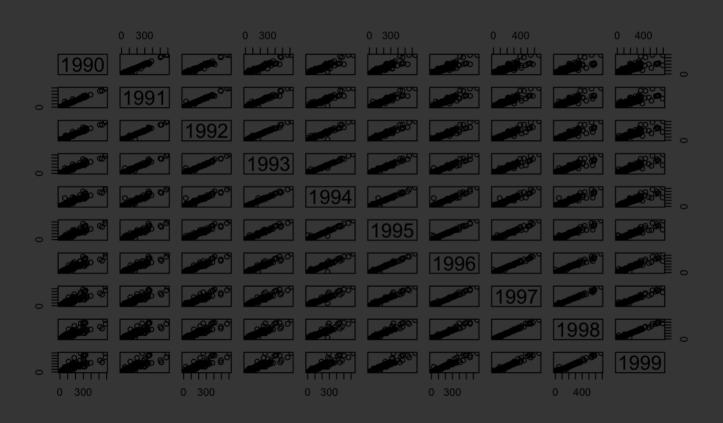
Matrix plot

library(GGally) ggpairs(avgs)



Matrix plot

pairs (avgs)



Lab Part 4

Website

Conclusion

- · Base R has apply statements that perform things repeatedly.
- · dplyr has a lot of more intuitive syntax.
 - group_by is very powerful, especilly with summarise/summarize
- Base R has good things for quickly summarizing rows or columnns of all numeric data.
 - The matrixStats package extends this to colMedians, colMaxs, etc.

Website

Website

Base R Plots - not covered

Data Summarization/Visualization

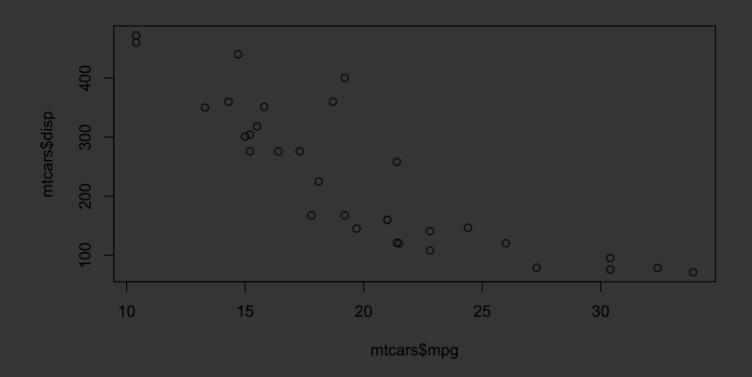
- Basic summarization plots
 - plot(x,y):scatterplot of x and y
 - boxplot (y~x): boxplot of y against levels of x
 - hist(x): histogram of x
 - density(x): kernel density plot of x

Basic Plots

Plotting is an important component of exploratory data analysis. We will review some of the more useful and informative plots here. We will go over formatting and making plots look nicer in additional lectures.

Scatterplot

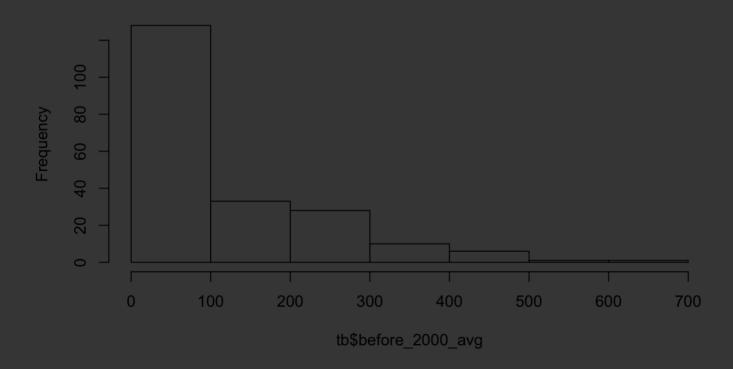
plot(mtcars\$mpg, mtcars\$disp)



Histograms

hist(tb\$before_2000_avg)

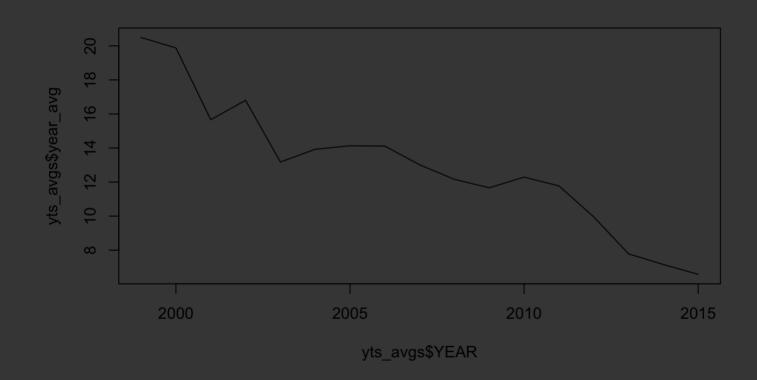
Histogram of tb\$before_2000_avg



Plot with a line

type = "1" means a line

plot(yts_avgs\$YEAR, yts_avgs\$year_avg, type = "1")

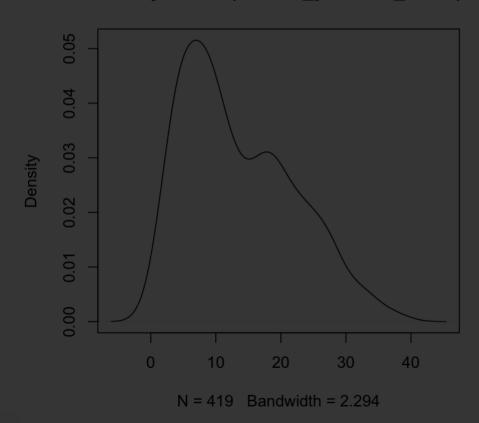


Density

Over all years and states, this is the density of smoking status incidence:

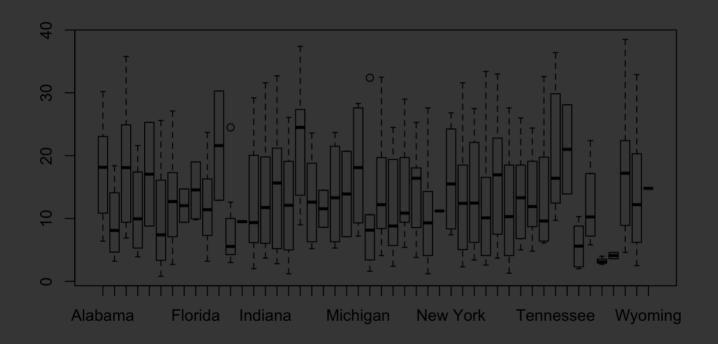
plot(density(sub_yts\$Data_Value))

density.default(x = sub_yts\$Data_Value)



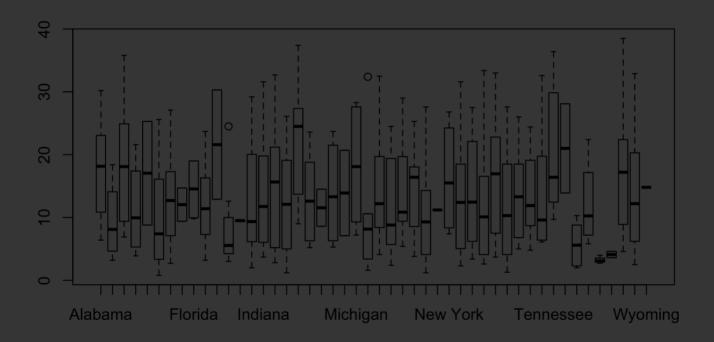
Boxplots

boxplot(sub yts\$Data Value ~ sub yts\$LocationDesc)



Boxplots

boxplot(Data Value ~ LocationDesc, data = sub yts)



Data Summarization for data.frames

- Basic summarization plots
 - matplot(x,y): scatterplot of two matrices, x and y
 - pairs (x, y): plots pairwise scatter plots of matrices x and y, column by column

Matrix plot

pairs (avgs)



Lab Part 4

Website