Data Summarization

Introduction to R for Public Health Researchers

Data Summarization

- Basic statistical summarization

    - `mean(x)`: takes the mean of x

    - `sd(x)`: takes the standard deviation of x

    - `median(x)`: takes the median of x

    - `quantile(x)`: displays sample quantities of x. Default is min, IQR, max

    - `range(x)`: displays the range. Same as c(min(x), max(x))

    - `sum(x)`: sum of x

- Transformations

    - `log` - log (base e) transformation

    - `log2` - log base 2 transform

    - `log10` - log base 10 transform

    - `sqrt` - square root

Some examples

We can use the `mtcars` to explore different ways of summarizing data. The `head` command displays the first `6` (default) rows of an object:

```
head(mtcars)
```

```
                   mpg cyl disp  hp drat    wt  qsec vs am gear carb
Mazda RX4         21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
Mazda RX4 Wag     21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
Datsun 710        22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
Hornet 4 Drive    21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
Valiant           18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

Statistical summarization

Note - the `$` references/selects columns from a `data.frame`/`tibble`:

```
mean(mtcars$hp)
```

```
[1] 146.6875
```

```
quantile(mtcars$hp)
```

```
    0%    25%    50%    75%   100%
  52.0   96.5  123.0  180.0  335.0
```

Statistical summarization

```
median(mtcars$wt)
```

```
[1] 3.325
```

```
quantile(mtcars$wt, probs = 0.6)
```

```
 60%
3.44
```

Statistical summarization

`t.test` will be covered more in detail later, gives a mean and 95% CI:

```
t.test(mtcars$wt)

	One Sample t-test

data:  mtcars$wt
t = 18.6, df = 31, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 2.864478 3.570022
sample estimates:
mean of x
  3.21725
```

Statistical summarization

Note that many of these functions have additional inputs regarding missing data, typically requiring the `na.rm` argument ("remove NAs").

```r
x = c(1,5,7,NA,4,2, 8,10,45,42)
mean(x)
```

```
[1] NA
```

```r
mean(x, na.rm = TRUE)
```

```
[1] 13.77778
```

```r
quantile(x, na.rm = TRUE)
```

```
  0%   25%   50%   75% 100%
   1     4     7    10    45
```

Data Summarization on matrices/data frames

- Basic statistical summarization

    - `rowMeans(x)`: takes the means of each row of x

    - `colMeans(x)`: takes the means of each column of x

    - `rowSums(x)`: takes the sum of each row of x

    - `colSums(x)`: takes the sum of each column of x

    - `summary(x)`: for data frames, displays the quantile information

Lab Part 1

Website

TB Incidence

Please download the TB incidence data:

http://johnmuschelli.com/intro_to_r/data/tb_incidence.xlsx

Here we will read in a `data.frame` of values from TB incidence:

```
library(readxl)
tb <- read_excel("../data/tb_incidence.xlsx")
head(tb)

# A tibble: 6 x 19
  `TB incidence, all forms (per 100 000 population per year)` `1990`
                                                       <chr>   <dbl>
1                                                Afghanistan     168
2                                                    Albania      25
3                                                    Algeria      38
4                                             American Samoa      21
5                                                    Andorra      36
6                                                     Angola     205
# ... with 17 more variables: `1991` <dbl>, `1992` <dbl>, `1993` <dbl>,
#   `1994` <dbl>, `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>,
#   `1999` <dbl>, `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>,
#   `2004` <dbl>, `2005` <dbl>, `2006` <dbl>, `2007` <dbl>
```

Indicator of TB

We can rename the first column to be the country measured using the `rename` function in `dplyr` (we have to use the ` things because there are spaces in the name):

```
library(dplyr)
tb = rename(tb,
        country = `TB incidence, all forms (per 100 000 population per yea
```

`colnames` will show us the column names and sho that country is renamed:

```
colnames(tb)
```

```
 [1] "country" "1990"    "1991"    "1992"    "1993"    "1994"    "1995"
 [8] "1996"    "1997"    "1998"    "1999"    "2000"    "2001"    "2002"
[15] "2003"    "2004"    "2005"    "2006"    "2007"
```

Column and Row means

`colMeans` and `rowMeans` must work on all numeric data. We will subset years before 2000 (starting with 1):

```
avgs = select(tb, starts_with("1"))
colMeans(avgs, na.rm = TRUE)
```

```
    1990     1991     1992     1993     1994     1995     1996     1997
105.5797 107.6715 108.3140 110.3188 111.9662 114.1981 115.3527 118.8792
    1998     1999
121.5169 125.0435
```

```
tb$before_2000_avg = rowMeans(avgs, na.rm = TRUE)
head(tb[, c("country", "before_2000_avg")])
```

```
# A tibble: 6 x 2
          country before_2000_avg
            <chr>           <dbl>
1     Afghanistan           168.0
2         Albania            26.3
3         Algeria            41.8
4  American Samoa             8.5
5         Andorra            28.8
6          Angola           224.6
```

Summary

Using `summary` can give you rough snapshots of each column, but you would likely use `mean`, `min`, `max`, and `quantile` when necessary:

```
summary(tb)
```

```
   country              1990             1991             1992
 Length:208        Min.   :  0.0    Min.   :  4.0    Min.   :  2.0
 Class :character  1st Qu.: 27.5    1st Qu.: 27.0    1st Qu.: 27.0
 Mode  :character  Median : 60.0    Median : 58.0    Median : 56.0
                   Mean   :105.6    Mean   :107.7    Mean   :108.3
                   3rd Qu.:165.0    3rd Qu.:171.0    3rd Qu.:171.5
                   Max.   :585.0    Max.   :594.0    Max.   :606.0
                   NA's   :1        NA's   :1        NA's   :1
      1993             1994             1995             1996
 Min.   :  4.0    Min.   :  0      Min.   :  3.0    Min.   :  0.0
 1st Qu.: 27.5    1st Qu.: 26      1st Qu.: 26.5    1st Qu.: 25.5
 Median : 56.0    Median : 57      Median : 58.0    Median : 60.0
 Mean   :110.3    Mean   :112      Mean   :114.2    Mean   :115.4
 3rd Qu.:171.0    3rd Qu.:174      3rd Qu.:177.5    3rd Qu.:179.0
 Max.   :618.0    Max.   :630      Max.   :642.0    Max.   :655.0
 NA's   :1        NA's   :1        NA's   :1        NA's   :1
      1997             1998             1999             2000
 Min.   :  0.0    Min.   :  0.0    Min.   :  0.0    Min.   :  0.0
 1st Qu.: 24.5    1st Qu.: 23.5    1st Qu.: 22.5    1st Qu.: 21.5
 Median : 64.0    Median : 63.0    Median : 66.0    Median : 60.0
 Mean   :118.9    Mean   :121.5    Mean   :125.0    Mean   :127.8
 3rd Qu.:181.0    3rd Qu.:188.5    3rd Qu.:192.5    3rd Qu.:191.0
```

Apply statements

You can apply more general functions to the rows or columns of a matrix or data frame, beyond the mean and sum.

```
apply(X, MARGIN, FUN, ...)
```

X : an array, including a matrix.

MARGIN : a vector giving the subscripts which the function will be applied over. E.g., for a matrix 1 indicates rows, 2 indicates columns, c(1, 2) indicates rows and columns. Where X has named dimnames, it can be a character vector selecting dimension names.

FUN : the function to be applied: see 'Details'.

… : optional arguments to FUN.

Apply statements

```r
apply(avgs,2,mean,na.rm=TRUE) # column means
```

```
     1990     1991     1992     1993     1994     1995     1996     1997
105.5797 107.6715 108.3140 110.3188 111.9662 114.1981 115.3527 118.8792
     1998     1999
121.5169 125.0435
```

```r
apply(avgs,2,sd,na.rm=TRUE) # columns sds
```

```
     1990     1991     1992     1993     1994     1995     1996     1997
110.6440 112.7687 114.4853 116.6744 120.0931 122.7119 126.1800 131.0858
     1998     1999
137.3754 146.0755
```

```r
apply(avgs,2,max,na.rm=TRUE) # column maxs
```

```
1990 1991 1992 1993 1994 1995 1996 1997 1998 1999
 585  594  606  618  630  642  655  668  681  695
```

Other Apply Statements

- `tapply()`: 'grouping' apply

- `lapply()`: 'list' apply [tomorrow]

- `sapply()`: 'simple' apply [tomorrow]

- Other less used ones…

See more details here: http://nsaunders.wordpress.com/2010/08/20/a-brief-introduction-to-apply-in-r/

Youth Tobacco Survey

Please download the Youth Tobacco Survey data. You can also read it in directly from the web:

```
library(readr)
yts = read_csv(
  "http://johnmuschelli.com/intro_to_r/data/Youth_Tobacco_Survey_YTS_Data.csv"
head(yts)
```

```
# A tibble: 6 x 31
   YEAR LocationAbbr LocationDesc                     TopicType
  <int>        <chr>        <chr>                          <chr>
1  2015           AZ      Arizona Tobacco Use - Survey Data
2  2015           AZ      Arizona Tobacco Use - Survey Data
3  2015           AZ      Arizona Tobacco Use - Survey Data
4  2015           AZ      Arizona Tobacco Use - Survey Data
5  2015           AZ      Arizona Tobacco Use - Survey Data
6  2015           AZ      Arizona Tobacco Use - Survey Data
# ... with 27 more variables: TopicDesc <chr>, MeasureDesc <chr>,
#   DataSource <chr>, Response <chr>, Data_Value_Unit <chr>,
#   Data_Value_Type <chr>, Data_Value <dbl>,
#   Data_Value_Footnote_Symbol <chr>, Data_Value_Footnote <chr>,
#   Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,
#   High_Confidence_Limit <dbl>, Sample_Size <int>, Gender <chr>,
#   Race <chr>, Age <chr>, Education <chr>, GeoLocation <chr>,
#   TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>,
#   StratificationID1 <chr>, StratificationID2 <chr>,
```

Length and unique

`unique(x)` will return the unique elements of `x`

```
unique(yts$LocationDesc)[1:10]
```

```
 [1] "Arizona"                "Connecticut"
 [3] "Georgia"                "Hawaii"
 [5] "Illinois"               "Louisiana"
 [7] "Mississippi"            "Utah"
 [9] "Missouri"               "National (States and DC)"
```

`length` will tell you the length of a vector. Combined with `unique`, tells you the number of unique elements:

```
length(unique(yts$LocationDesc))
```

```
[1] 50
```

Table

`table(x)` will return a frequency table of unique elements of `x`

```
table(yts$LocationDesc)[1:5]
```

Alabama	Arizona	Arkansas	California	Colorado
378	240	210	96	48

Lab Part 2

Website

Subsetting to specific columns

Let's just take smoking status measures for all genders using `filter`, and the columns that represent the year, state using `select`:

```r
library(dplyr)
sub_yts = filter(yts,
                 MeasureDesc == "Smoking Status",
                 Gender == "Overall",
                 Response == "Current")
sub_yts = select(sub_yts, YEAR, LocationDesc, Data_Value)
head(sub_yts, 4)
```

```
# A tibble: 4 x 3
   YEAR LocationDesc Data_Value
  <int>        <chr>      <dbl>
1  2015      Arizona        3.2
2  2015  Connecticut        0.8
3  2015  Connecticut        5.6
4  2015      Georgia       10.8
```

Perform Operations By Groups: dplyr

`group_by` allows you group the data in a more intuitive way than `tapply`

We will use `group_by` to group the data by line, then use `summarize` (or `summarise`) to get the mean percentage of current smokers:

```
summarize(group_by(sub_yts, YEAR), year_avg = mean(Data_Value, na.rm = TRUE))

# A tibble: 17 x 2
     YEAR  year_avg
    <int>     <dbl>
 1   1999 20.493333
 2   2000 19.878431
 3   2001 15.661111
 4   2002 16.802326
 5   2003 13.176190
 6   2004 13.926923
 7   2005 14.128571
 8   2006 14.113636
 9   2007 13.013636
10   2008 12.159091
11   2009 11.663333
12   2010 12.290000
13   2011 11.773913
14   2012  9.954545
15   2013  7.782759
```

Using the `pipe` (comes with `dplyr`):

Pipe `sub_yts` into `group_by`, then pipe that into `summarize`:

```
yts_avgs = sub_yts %>%
  group_by(YEAR) %>%
  summarize(year_avg = mean(Data_Value, na.rm = TRUE))
head(yts_avgs)

# A tibble: 6 x 2
    YEAR year_avg
   <int>    <dbl>
1   1999 20.49333
2   2000 19.87843
3   2001 15.66111
4   2002 16.80233
5   2003 13.17619
6   2004 13.92692
```

Counting

Standard statistics can be calculated. There are other functions, such as `n()` count the number of observations, `tally()` to count as a wrapper:

```
sub_yts %>%
  group_by(YEAR) %>%
  summarize(n = n()) %>%
  head
```

```
# A tibble: 6 x 2
   YEAR      n
  <int> <int>
1  1999     15
2  2000     51
3  2001     18
4  2002     43
5  2003     21
6  2004     26
```

```
sub_yts %>%
  group_by(YEAR) %>%
  tally() %>%
  head
```

```
# A tibble: 6 x 2
   YEAR      n
  <int> <int>
```

Lab Part 3

Website

Data Summarization/Visualization: ggplot2

`ggplot2` is a package of plotting that is very popular and powerful (using the **g**rammar of **g**raphics). We will use `qplot` ("quick plot") for most of the basic examples:

```
qplot

function (x, y = NULL, ..., data, facets = NULL, margins = FALSE,
    geom = "auto", xlim = c(NA, NA), ylim = c(NA, NA), log = "",
    main = NULL, xlab = deparse(substitute(x)), ylab = deparse(substitute(y)),
    asp = NA, stat = NULL, position = NULL)
{
    if (!missing(stat))
        warning("`stat` is deprecated", call. = FALSE)
    if (!missing(position))
        warning("`position` is deprecated", call. = FALSE)
    if (!is.character(geom))
        stop("`geom` must be a character vector", call. = FALSE)
    argnames <- names(as.list(match.call(expand.dots = FALSE)[-1]))
    arguments <- as.list(match.call()[-1])
    env <- parent.frame()
    aesthetics <- compact(arguments[.all_aesthetics])
    aesthetics <- aesthetics[!is.constant(aesthetics)]
    aes_names <- names(aesthetics)
    aesthetics <- rename_aes(aesthetics)
    class(aesthetics) <- "uneval"
    if (missing(data)) {
```

Basic Plots

Plotting is an important component of exploratory data analysis. We will review some of the more useful and informative plots here. We will go over formatting and making plots look nicer in additional lectures.

Scatterplot

```
library(ggplot2)
qplot(x = disp, y = mpg, data = mtcars)
```

Histograms

```
qplot(x = before_2000_avg, data = tb, geom = "histogram")
```

```
Warning: Removed 1 rows containing non-finite values (stat_bin).
```
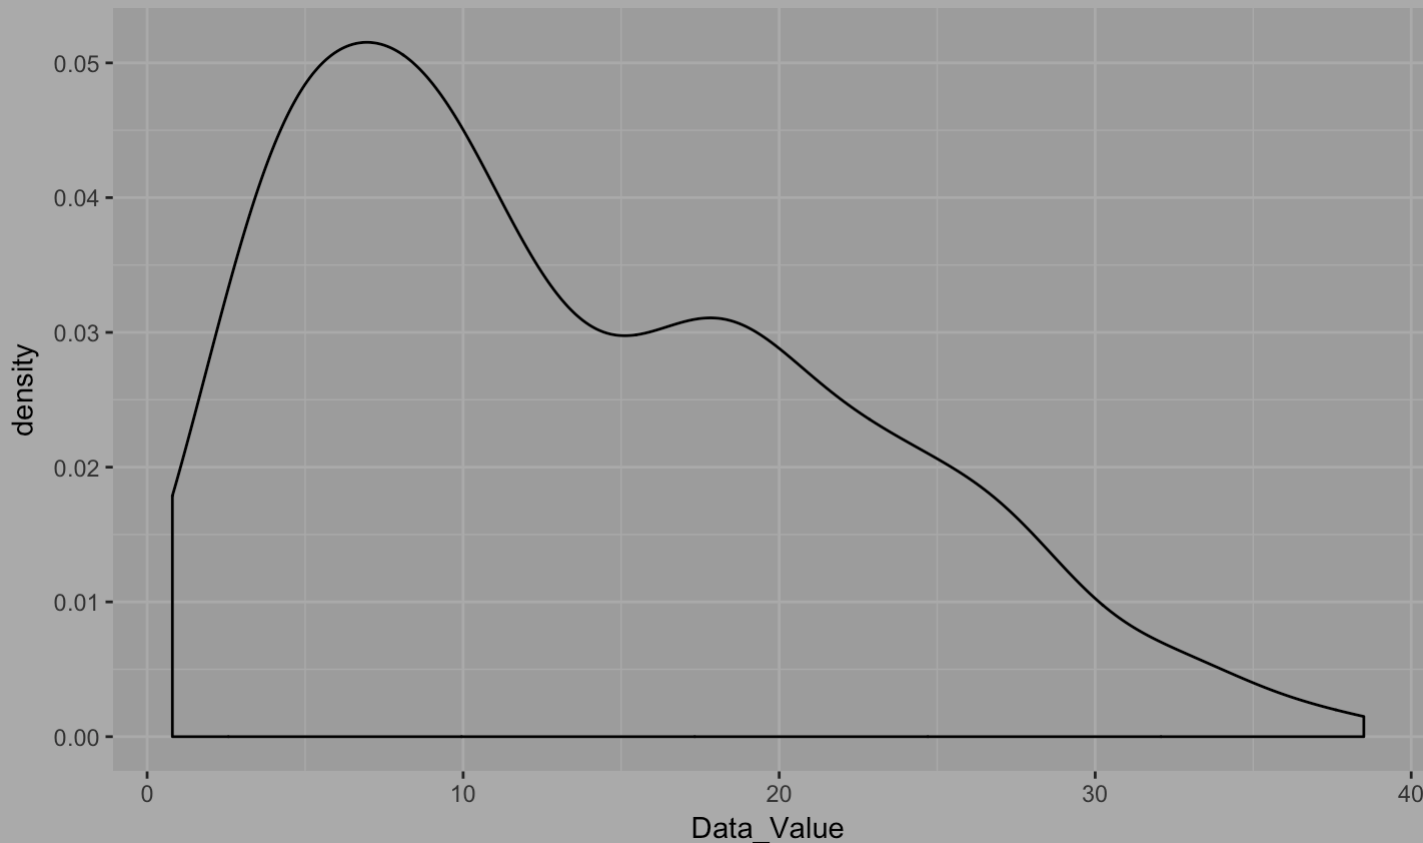
Plot with a line

```
qplot(x = YEAR, y = year_avg, data = yts_avgs, geom = "line")
```

Density

Over all years and states, this is the density of smoking status incidence:

```
qplot(x = Data_Value, data = sub_yts, geom = "density")
```


Boxplots

```
qplot(x = LocationDesc, y = Data_Value, data = sub_yts, geom = "boxplot")
```

Data Summarization for data.frames

- Basic summarization plots

  - `matplot(x,y)`: scatterplot of two matrices, x and y

  - `pairs(x,y)`: plots pairwise scatter plots of matrices x and y, column by column

Matrix plot

```
pairs(avgs)
```
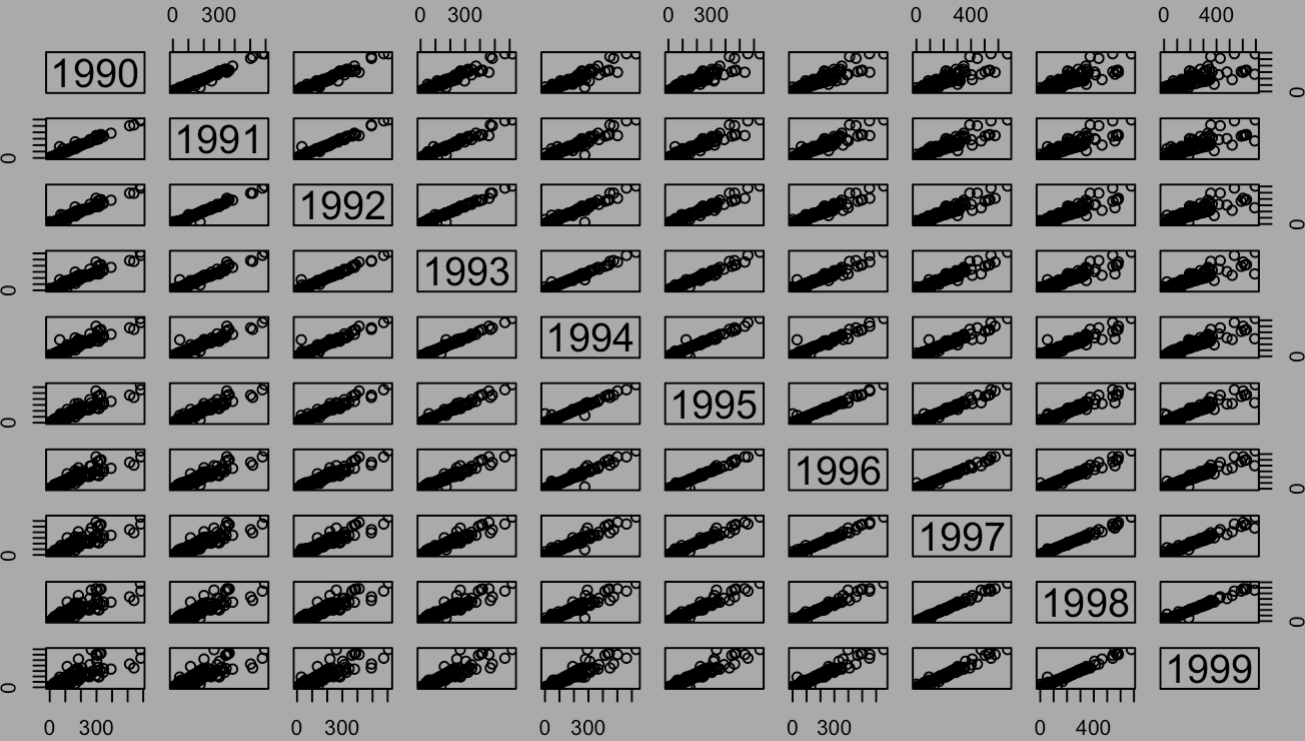
Lab Part 4

Website

Conclusion

- Base R has apply statements that perform things repeatedly.

- `dplyr` has a lot of more intuitive syntax.

    - `group_by` is very powerful, especilly with `summarise/summarize`

- Base R has good things for quickly summarizing rows or colummns of all numeric data.

    - The `matrixStats` package extends this to `colMedians, colMaxs,` etc.

Website

Website

Base R Plots - not covered

Data Summarization/Visualization

- Basic summarization plots
    - `plot(x,y)`: scatterplot of x and y

    - `boxplot(y~x)`: boxplot of y against levels of x

    - `hist(x)`: histogram of x

    - `density(x)`: kernel density plot of x

Basic Plots

Plotting is an important component of exploratory data analysis. We will review some of the more useful and informative plots here. We will go over formatting and making plots look nicer in additional lectures.

Scatterplot

```
plot(mtcars$mpg, mtcars$disp)
```

Histograms

```
hist(tb$before_2000_avg)
```

Histogram of tb$before_2000_avg
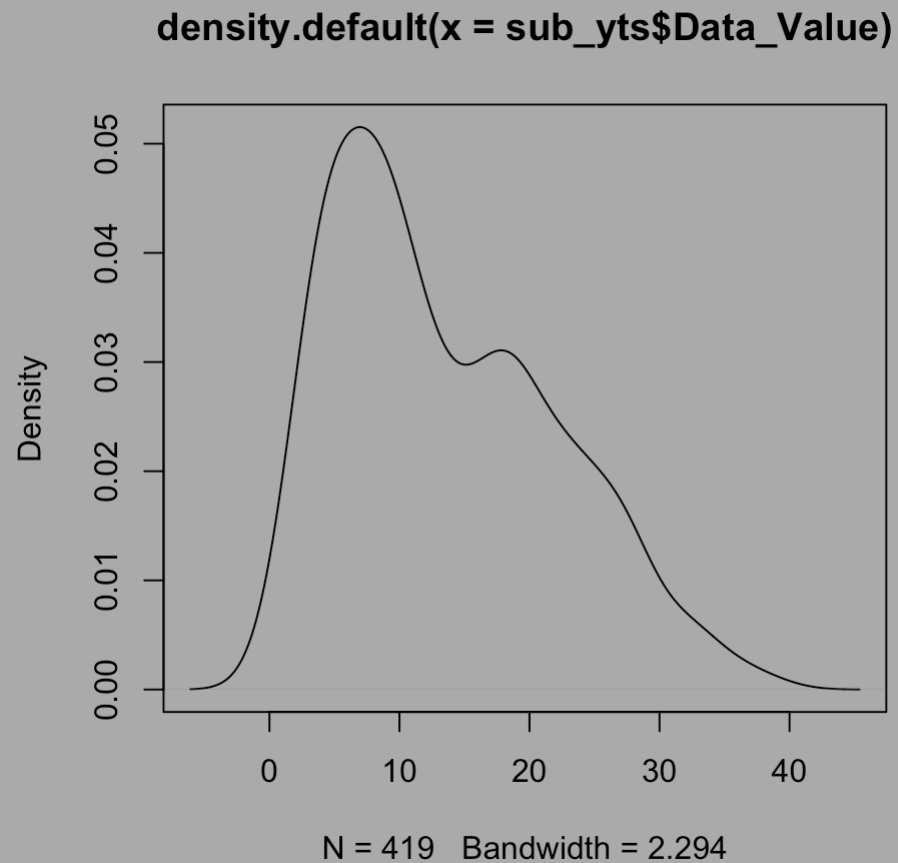
Plot with a line

`type = "l"` means a line

```
plot(yts_avgs$YEAR, yts_avgs$year_avg, type = "l")
```

Density

Over all years and states, this is the density of smoking status incidence:

```
plot(density(sub_yts$Data_Value))
```
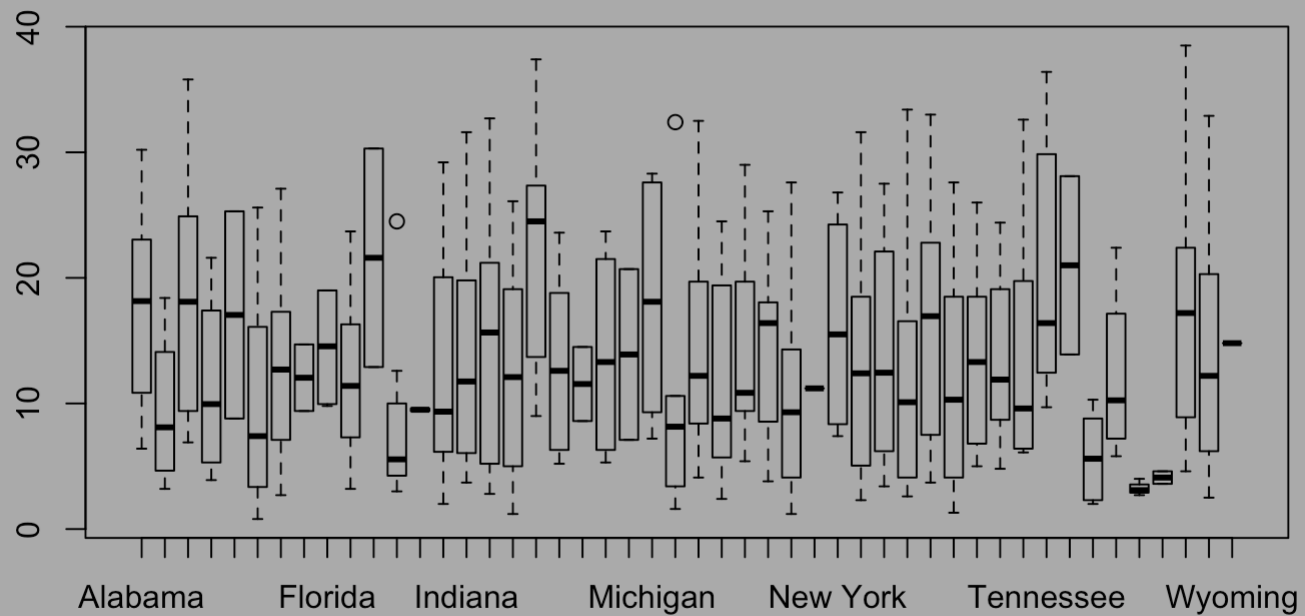

density.default(x = sub_yts$Data_Value)

N = 419   Bandwidth = 2.294

Boxplots

```
boxplot(sub_yts$Data_Value ~ sub_yts$LocationDesc)
```

Boxplots

```
boxplot(Data_Value ~ LocationDesc, data = sub_yts)
```

Data Summarization for data.frames

- Basic summarization plots

    - `matplot(x,y)`: scatterplot of two matrices, x and y

    - `pairs(x,y)`: plots pairwise scatter plots of matrices x and y, column by column

Matrix plot

```
pairs(avgs)
```

Lab Part 4

Website