# Manipulating Data in R

Introduction to R for Public Health Researchers

# Reshaping Data

In this module, we will show you how to:

1. Reshaping data from wide (fat) to long (tall)

2. Reshaping data from long (tall) to wide (fat)

3. Merging Data

4. Perform operations by a grouping variable

# Setup

We will show you how to do each operation in base R then show you how to use the `dplyr` or `tidyr` package to do the same operation (if applicable).

See the "Data Wrangling Cheat Sheet using `dplyr` and `tidyr`":

- https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf

# What is wide/long data?

See http://www.cookbook-r.com/Manipulating_data/Converting_data_between_wide_and_long_format/

- Wide - multiple columns per observation

  - e.g. visit1, visit2, visit3

```
  id visit1 visit2 visit3
1  1     10      4      3
2  2      5      6     NA
```

- Long - multiple rows per observation

```
  id visit value
1  1     1    10
2  1     2     4
3  1     3     3
4  2     1     5
5  2     2     6
```

More accurately, data is wide or long **with respect** to certain variables.

# Data used: Charm City Circulator

http://johnmuschelli.com/intro_to_r/data/Charm_City_Circulator_Ridership.csv

```r
circ = read_csv(
  paste0("http://johnmuschelli.com/intro_to_r/",
         "data/Charm_City_Circulator_Ridership.csv")
)
head(circ, 2)
```

```
# A tibble: 2 x 15
      day         date orangeBoardings orangeAlightings orangeAverage
    <chr>        <chr>           <int>            <int>         <dbl>
1  Monday 01/11/2010             877             1027           952
2 Tuesday 01/12/2010             777              815           796
# ... with 10 more variables: purpleBoardings <int>,
#   purpleAlightings <int>, purpleAverage <dbl>, greenBoardings <int>,
#   greenAlightings <int>, greenAverage <dbl>, bannerBoardings <int>,
#   bannerAlightings <int>, bannerAverage <dbl>, daily <dbl>
```

# Creating a Date class from a character date

```r
library(lubridate) # great for dates!
library(dplyr)  # mutate/summarise functions
```

# Creating a Date class from a character date

```
sum(is.na(circ$date))
```

```
[1] 0
```

```
sum( circ$date == "")
```

```
[1] 0
```

```
circ = mutate(circ, date = mdy(date))
sum( is.na(circ$date) ) # all converted correctly
```

```
[1] 0
```

```
head(circ$date, 3)
```

```
[1] "2010-01-11" "2010-01-12" "2010-01-13"
```

```
class(circ$date)
```

```
[1] "Date"
```

# Reshaping data from wide (fat) to long (tall): base R

The `reshape` command exists. It is a **confusing** function. Don't use it.

# tidyr package

`tidyr` allows you to "tidy" your data. We will be talking about:

- `gather` - make multiple columns into variables, (wide to long)
- `spread` - make a variable into multiple columns, (long to wide)
- `separate` - string into multiple columns
- `unite` - multiple columns into one string

# Reshaping data from wide (fat) to long (tall): tidyr

`tidyr::gather` - puts column data into rows.

We want the column names into "`var`" variable in the output dataset and the value in "`number`" variable. We then describe which columns we want to "gather:"

```r
library(tidyr)
long = gather(circ, key = "var", value = "number",
              starts_with("orange"), starts_with("purple"),
              starts_with("green"), starts_with("banner"))
head(long, 4)
```

```
# A tibble: 4 x 5
        day       date  daily               var number
      <chr>     <date>  <dbl>             <chr>  <dbl>
1    Monday 2010-01-11  952.0 orangeBoardings    877
2   Tuesday 2010-01-12  796.0 orangeBoardings    777
3 Wednesday 2010-01-13 1211.5 orangeBoardings   1203
4  Thursday 2010-01-14 1213.5 orangeBoardings   1194
```

# Reshaping data from wide (fat) to long (tall): tidyr

```
table(long$var)
```

```
    bannerAlightings       bannerAverage    bannerBoardings    greenAlightings
                1146                1146               1146               1146
        greenAverage       greenBoardings   orangeAlightings     orangeAverage
                1146                1146               1146               1146
     orangeBoardings     purpleAlightings      purpleAverage    purpleBoardings
                1146                1146               1146               1146
```

# Making a separator

We will use `str_replace` from `stringr` to put periods in the names:

```r
library(stringr)
long = long %>% mutate(
  var = var %>% str_replace("Board", ".Board") %>%
    str_replace("Alight", ".Alight") %>%
    str_replace("Average", ".Average")
)
table(long$var)
```

```
 banner.Alightings      banner.Average    banner.Boardings    green.Alightings
              1146                1146                1146                1146
     green.Average     green.Boardings  orange.Alightings     orange.Average
              1146                1146                1146                1146
  orange.Boardings   purple.Alightings     purple.Average    purple.Boardings
              1146                1146                1146                1146
```

# Reshaping data from wide (fat) to long (tall): tidyr

Now each `var` is boardings, averages, or alightings. We want to separate these so we can have these by line. Remember "." is special character:

```
long = separate(long, var, into = c("line", "type"),
                sep = "[.]")
head(long, 2)
```

```
# A tibble: 2 x 6
     day       date daily   line      type number
   <chr>     <date> <dbl>  <chr>     <chr>  <dbl>
1  Monday 2010-01-11   952 orange Boardings    877
2 Tuesday 2010-01-12   796 orange Boardings    777
```

```
unique(long$line)
```

```
[1] "orange" "purple" "green"  "banner"
```

```
unique(long$type)
```

```
[1] "Boardings"  "Alightings" "Average"
```

# Re-uniting all the lines

If we had the opposite problem, we could use the `unite` function:

```
reunited = long %>%
  unite(col = var, line, type, sep = ".")
reunited %>% select(day, var) %>% head(3) %>% print

# A tibble: 3 x 2
        day               var
      <chr>             <chr>
1    Monday  orange.Boardings
2   Tuesday  orange.Boardings
3 Wednesday  orange.Boardings
```

We could also use `paste/paste0`.

# Making column names a little more separated

Alternative: We could have replaced the column names first **then** reshaped:

```
cn = colnames(circ)
cn = cn %>%
  str_replace("Board", ".Board") %>%
  str_replace("Alight", ".Alight") %>%
  str_replace("Average", ".Average")
colnames(circ) = cn # then reshape using gather!
```

# Reshaping data from long (tall) to wide (fat): tidyr

In `tidyr`, the `spread` function spreads rows into columns. Now we have a long data set, but we want to separate the Average, Alightings and Boardings into different columns:

```
# have to remove missing days
wide = filter(long, !is.na(date))
wide = spread(wide, type, number)
head(wide)
```

```
# A tibble: 6 x 7
      day        date   daily    line Alightings Average Boardings
    <chr>      <date>   <dbl>   <chr>      <dbl>   <dbl>     <dbl>
1 Friday 2010-01-15  1644.0  banner         NA      NA        NA
2 Friday 2010-01-15  1644.0   green         NA      NA        NA
3 Friday 2010-01-15  1644.0  orange       1643    1644      1645
4 Friday 2010-01-15  1644.0  purple         NA      NA        NA
5 Friday 2010-01-22  1394.5  banner         NA      NA        NA
6 Friday 2010-01-22  1394.5   green         NA      NA        NA
```

# Reshaping data from long (tall) to wide (fat): tidyr

We can use `rowSums` to see if any values in the row is `NA` and keep if the row, which is a combination of date and line type has any non-missing data.

```r
# wide = wide %>%
#     select(Alightings, Average, Boardings) %>%
#     mutate(good = rowSums(is.na(.)) > 0)
not_namat = !is.na(select(wide, Alightings, Average, Boardings))
head(not_namat, 2)
```

```
  Alightings Average Boardings
1      FALSE   FALSE     FALSE
2      FALSE   FALSE     FALSE
```

```r
wide$good = rowSums(not_namat) > 0
```

# Reshaping data from long (tall) to wide (fat): tidyr

Now we can filter only the good rows and delete the `good` column.

```
wide = filter(wide, good) %>% select(-good)
head(wide)

# A tibble: 6 x 7
     day        date   daily    line Alightings Average Boardings
   <chr>      <date>   <dbl>   <chr>      <dbl>   <dbl>     <dbl>
1 Friday 2010-01-15  1644.0  orange       1643  1644.0      1645
2 Friday 2010-01-22  1394.5  orange       1388  1394.5      1401
3 Friday 2010-01-29  1332.0  orange       1322  1332.0      1342
4 Friday 2010-02-05  1217.5  orange       1204  1217.5      1231
5 Friday 2010-02-12   671.0  orange        678   671.0       664
6 Friday 2010-02-19  1642.0  orange       1647  1642.0      1637
```

# Finding the First (or Last) record

```r
long = long %>% filter(!is.na(number) & number > 0)
first_and_last = long %>% arrange(date) %>% # arrange by date
  filter(type %in% "Boardings") %>% # keep boardings only
  group_by(line) %>% # group by line
  slice( c(1, n())) # select ("slice") first and last (n() command) lines
first_and_last %>%  head(4)
```

```
# A tibble: 4 x 6
# Groups:   line [2]
      day        date   daily   line       type number
    <chr>      <date>   <dbl>  <chr>      <chr>  <dbl>
1  Monday 2012-06-04 13342.5 banner  Boardings    520
2  Friday 2013-03-01      NA banner  Boardings    817
3 Tuesday 2011-11-01  8873.0  green  Boardings    887
4  Friday 2013-03-01      NA  green  Boardings   2592
```

# Data Merging/Append in Base R

- Merging - joining data sets together - usually on key variables, usually "id"
- `merge()` is the most common way to do this with data sets
- `rbind`/`cbind` - row/column bind, respectively
    - `rbind` is the equivalent of "appending" in Stata or "setting" in SAS
    - `cbind` allows you to add columns in addition to the previous ways
- `t()` is a function that will transpose the data

# Merging

```
base <- data.frame(id = 1:10, Age= seq(55,60, length=10))
base[1:2,]
```

```
   id       Age
1   1 55.00000
2   2 55.55556
```

```
visits <- data.frame(id = rep(1:8, 3), visit= rep(1:3, 8),
                     Outcome = seq(10,50, length=24))
visits[1:2,]
```

```
  id visit  Outcome
1  1     1 10.00000
2  2     2 11.73913
```

# Merging

```
merged.data <- merge(base, visits, by="id")
merged.data[1:5,]
```

```
  id      Age visit  Outcome
1  1 55.00000     1 10.00000
2  1 55.00000     3 23.91304
3  1 55.00000     2 37.82609
4  2 55.55556     2 11.73913
5  2 55.55556     1 25.65217
```

```
dim(merged.data)
```

```
[1] 24  4
```

# Merging

```
all.data <- merge(base, visits, by="id", all=TRUE)
tail(all.data)
```

```
   id      Age visit  Outcome
21  7 58.33333     2 48.26087
22  8 58.88889     2 22.17391
23  8 58.88889     1 36.08696
24  8 58.88889     3 50.00000
25  9 59.44444    NA       NA
26 10 60.00000    NA       NA
```

```
dim(all.data)
```

```
[1] 26  4
```

# Joining in `dplyr`

- `?join` - see different types of joining for `dplyr`

- Let's look at https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf

# Left Join

```
lj = left_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(lj)
```

```
[1] 26  4
```

```
tail(lj)
```

```
    id      Age visit  Outcome
21   7 58.33333     2 48.26087
22   8 58.88889     2 22.17391
23   8 58.88889     1 36.08696
24   8 58.88889     3 50.00000
25   9 59.44444    NA       NA
26  10 60.00000    NA       NA
```

# Right Join

```
rj = right_join(base, visits)

Joining, by = "id"

dim(rj)

[1] 24  4

tail(rj)
```

```
   id      Age visit  Outcome
19  3 56.11111     1 41.30435
20  4 56.66667     2 43.04348
21  5 57.22222     3 44.78261
22  6 57.77778     1 46.52174
23  7 58.33333     2 48.26087
24  8 58.88889     3 50.00000
```

# Full Join

```
fj = full_join(base, visits)

Joining, by = "id"

dim(fj)

[1] 26  4

tail(fj)

   id      Age visit  Outcome
21  7 58.33333     2 48.26087
22  8 58.88889     2 22.17391
23  8 58.88889     1 36.08696
24  8 58.88889     3 50.00000
25  9 59.44444    NA       NA
26 10 60.00000    NA       NA
```

# Website

[Website](Website)