

Subsetting Data in R

Introduction to R for Public Health Researchers

Overview

We showed one way to read data into R using `read_csv` and `read.csv`. In this module, we will show you how to:

1. Select specific elements of an object by an index or logical condition
2. Renaming columns of a `data.frame`
3. Subset rows of a `data.frame`
4. Subset columns of a `data.frame`
5. Add/remove new columns to a `data.frame`
6. Order the columns of a `data.frame`
7. Order the rows of a `data.frame`

Setup

We will show you how to do each operation in base R then show you how to use the `dplyr` package to do the same operation (if applicable).

Many resources on how to use `dplyr` exist and are straightforward:

- <https://cran.rstudio.com/web/packages/dplyr/vignettes/>
- https://stat545-ubc.github.io/block009_dplyr-intro.html
- <https://www.datacamp.com/courses/dplyr-data-manipulation-r-tutorial>

The `dplyr` package also interfaces well with tibbles.

Select specific elements using an index

Often you only want to look at subsets of a data set at any given time. As a review, elements of an R object are selected using the brackets ([and]).

For example, `x` is a vector of numbers and we can select the second element of `x` using the brackets and an index (2):

```
x = c(1, 4, 2, 8, 10)
x[2]
```

```
[1] 4
```

Select specific elements using an index

We can select the fifth or second AND fifth elements below:

```
x = c(1, 2, 4, 8, 10)  
x[5]
```

```
[1] 10
```

```
x[c(2, 5)]
```

```
[1] 2 10
```

Subsetting by deletion of entries

You can put a minus (-) before integers inside brackets to remove these indices from the data.

```
x[-2] # all but the second
```

```
[1] 1 4 8 10
```

Note that you have to be careful with this syntax when dropping more than 1 element:

```
x[-c(1,2,3)] # drop first 3
```

```
[1] 8 10
```

```
# x[-1:3] # shorthand. R sees as -1 to 3  
x[-(1:3)] # needs parentheses
```

```
[1] 8 10
```

Select specific elements using logical operators

What about selecting rows based on the values of two variables? We use logical statements. Here we select only elements of `x` greater than 2:

```
x
```

```
[1] 1 2 4 8 10
```

```
x > 2
```

```
[1] FALSE FALSE  TRUE  TRUE  TRUE
```

```
x[ x > 2 ]
```

```
[1] 4 8 10
```

Select specific elements using logical operators

You can have multiple logical conditions using the following:

- `&` : AND
- `|` : OR

```
x[ x > 2 & x < 5 ]
```

```
[1] 4
```

```
x[ x > 5 | x == 2 ]
```

```
[1] 2 8 10
```


which function

The `which` functions takes in logical vectors and returns the index for the elements where the logical value is `TRUE`.

```
which(x > 5 | x == 2) # returns index
```

```
[1] 2 4 5
```

```
x[ which(x > 5 | x == 2) ]
```

```
[1] 2 8 10
```

```
x[ x > 5 | x == 2 ]
```

```
[1] 2 8 10
```

Creating a `data.frame` to work with

Here we use one of the datasets that comes with R called `mtcars` create a toy `data.frame` named `df` using random data:

```
data(mtcars)
df = mtcars
tbl = as.tbl(df)
```

Renaming Columns

Renaming Columns of a `data.frame`: base R

We can use the `colnames` function to directly reassign column names of `df`:

```
colnames(df)[1:3] = c("MPG", "CYL", "DISP")  
head(df)
```

	MPG	CYL	DISP	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
colnames(df)[1:3] = c("mpg", "cyl", "disp") #reset
```

Renaming Columns of a `data.frame`: base R

We can assign the column names, change the ones we want, and then re-assign the column names:

```
cn = colnames(df)
cn[ cn == "drat" ] = "DRAT"
colnames(df) = cn
head(df)
```

	mpg	cyl	disp	hp	DRAT	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
colnames(df)[ colnames(df) == "DRAT" ] = "drat" #reset
```

Renaming Columns of a `data.frame`: `dplyr`

```
library(dplyr)
```

Note, when loading `dplyr`, it says objects can be "masked". That means if you use a function defined in 2 places, it uses the one that is loaded in **last**.

Renaming Columns of a `data.frame`: `dplyr`

For example, if we print `filter`, then we see at the bottom `namespace:dplyr`, which means when you type `filter`, it will use the one from the `dplyr` package.

```
filter
```

```
function (.data, ...)  
{  
  filter_(.data, .dots = lazyeval::lazy_dots(...))  
}  
<environment: namespace:dplyr>
```

Renaming Columns of a `data.frame`: `dplyr`

A `filter` function exists by default in the `stats` package, however. If you want to make sure you use that one, you use `PackageName::Function` with the colon-colon ("`::`") operator.

```
head(stats::filter, 2)
```

```
1 function (x, filter, method = c("convolution", "recursive"),  
2       sides = 2L, circular = FALSE, init = NULL)
```

This is important when loading many packages, and you may have some conflicts/masking:

Renaming Columns of a `data.frame`: `dplyr`

To rename columns in `dplyr`, you use the `rename` command

```
df = dplyr::rename(df, MPG = mpg)
head(df)
```

	MPG	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
df = dplyr::rename(df, mpg = MPG) # reset
```

Subsetting Columns

Subset columns of a `data.frame`:

We can grab the `carb` column using the `$` operator.

```
df$carb
```

```
[1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2
```

Subset columns of a `data.frame`:

We can also subset a `data.frame` using the bracket `[,]` subsetting.

For `data.frames` and matrices (2-dimensional objects), the brackets are `[rows, columns]` subsetting. We can grab the `x` column using the index of the column or the column name (`"carb"`)

```
df[, 11]
```

```
[1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2
```

```
df[, "carb"]
```

```
[1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2
```

Biggest difference between `tbl` and `data.frame`:

Mostly, `tbl` (tibbles) are the same as `data.frame`s, except they don't print all lines. When subsetting only one column using brackets, a `data.frame` will return a vector, but a `tbl` will return a `tbl`

```
df[, 1]
```

```
[1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2
[15] 10.4 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4
[29] 15.8 19.7 15.0 21.4
```

```
tbl[, 1]
```

```
# A tibble: 32 x 1
```

```
  mpg
<dbl>
1  21.0
2  21.0
3  22.8
4  21.4
5  18.7
6  18.1
7  14.3
8  24.4
9  22.8
```

Subset columns of a `data.frame`:

We can select multiple columns using multiple column names:

```
df[, c("mpg", "cyl")]
```

	mpg	cyl
Mazda RX4	21.0	6
Mazda RX4 Wag	21.0	6
Datsun 710	22.8	4
Hornet 4 Drive	21.4	6
Hornet Sportabout	18.7	8
Valiant	18.1	6
Duster 360	14.3	8
Merc 240D	24.4	4
Merc 230	22.8	4
Merc 280	19.2	6
Merc 280C	17.8	6
Merc 450SE	16.4	8
Merc 450SL	17.3	8
Merc 450SLC	15.2	8
Cadillac Fleetwood	10.4	8
Lincoln Continental	10.4	8
Chrysler Imperial	14.7	8
Fiat 128	32.4	4
Honda Civic	30.4	4
Toyota Corolla	33.9	4

Subset columns of a `data.frame`: `dplyr`

The `select` command from `dplyr` allows you to subset

```
select(df, mpg)
```

	mpg
Mazda RX4	21.0
Mazda RX4 Wag	21.0
Datsun 710	22.8
Hornet 4 Drive	21.4
Hornet Sportabout	18.7
Valiant	18.1
Duster 360	14.3
Merc 240D	24.4
Merc 230	22.8
Merc 280	19.2
Merc 280C	17.8
Merc 450SE	16.4
Merc 450SL	17.3
Merc 450SLC	15.2
Cadillac Fleetwood	10.4
Lincoln Continental	10.4
Chrysler Imperial	14.7
Fiat 128	32.4
Honda Civic	30.4
Toyota Corolla	33.9

Select columns of a `data.frame`: `dplyr`

The `select` command from `dplyr` allows you to subset columns of

```
select(df, mpg, cyl)
```

	mpg	cyl
Mazda RX4	21.0	6
Mazda RX4 Wag	21.0	6
Datsun 710	22.8	4
Hornet 4 Drive	21.4	6
Hornet Sportabout	18.7	8
Valiant	18.1	6
Duster 360	14.3	8
Merc 240D	24.4	4
Merc 230	22.8	4
Merc 280	19.2	6
Merc 280C	17.8	6
Merc 450SE	16.4	8
Merc 450SL	17.3	8
Merc 450SLC	15.2	8
Cadillac Fleetwood	10.4	8
Lincoln Continental	10.4	8
Chrysler Imperial	14.7	8
Fiat 128	32.4	4
Honda Civic	30.4	4
Toyota Corolla	33.9	4

Subsetting Rows

Subset rows of a `data.frame` with indices:

Let's select **rows** 1 and 3 from `df` using brackets:

```
df[ c(1, 3), ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.62	16.46	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.32	18.61	1	1	4	1

Subset rows of a data.frame:

Let's select the rows of `df` where the `mpg` column is greater than 20 or is less than 14. Without any index for columns, all columns are returned:

```
df[ df$mpg > 20 | df$mpg < 14, ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Subset rows of a `data.frame`:

We can subset both rows and columns at the same time:

```
df[ df$mpg > 20 | df$mpg < 14, c("cyl", "hp") ]
```

	cyl	hp
Mazda RX4	6	110
Mazda RX4 Wag	6	110
Datsun 710	4	93
Hornet 4 Drive	6	110
Merc 240D	4	62
Merc 230	4	95
Cadillac Fleetwood	8	205
Lincoln Continental	8	215
Fiat 128	4	66
Honda Civic	4	52
Toyota Corolla	4	65
Toyota Corona	4	97
Camaro Z28	8	245
Fiat X1-9	4	66
Porsche 914-2	4	91
Lotus Europa	4	113
Volvo 142E	4	109

Subset rows of a `data.frame`: `dplyr`

The command in `dplyr` for subsetting rows is `filter`. Try `?filter`

```
filter(df, mpg > 20 | mpg < 14)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
6	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
7	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
8	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
9	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
10	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
11	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
12	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
13	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
14	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
15	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
16	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
17	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Note, no `$` or subsetting is necessary. R "knows" `mpg` refers to a column of `df`.

Subset rows of a `data.frame`: `dplyr`

By default, you can separate conditions by commas, and `filter` assumes these statements are joined by `&`

```
filter(df, mpg > 20 & cyl == 4)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
2	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
3	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
4	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
5	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
6	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
7	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
8	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
9	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
10	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
11	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

```
filter(df, mpg > 20, cyl == 4)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
2	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
3	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2

Combining `filter` and `select`

You can combine `filter` and `select` to subset the rows and columns, respectively, of a `data.frame`:

```
select(filter(df, mpg > 20 & cyl == 4), cyl, hp)
```

	cyl	hp
1	4	93
2	4	62
3	4	95
4	4	66
5	4	52
6	4	65
7	4	97
8	4	66
9	4	91
10	4	113
11	4	109

In \mathbb{R} , the common way to perform multiple operations is to wrap functions around each other in a nested way such as above

Assigning Temporary Objects

One can also create temporary objects and reassign them:

```
df2 = filter(df, mpg > 20 & cyl == 4)  
df2 = select(df2, cyl, hp)
```


Using the `pipe` (comes with `dplyr`):

Recently, the pipe `%>%` makes things such as this much more readable. It reads left side "pipes" into right side. RStudio CMD/Ctrl + Shift + M shortcut. Pipe `df` into `filter`, then pipe that into `select`:

```
df %>% filter(mpg > 20 & cyl == 4) %>% select(cyl, hp)
```

	cyl	hp
1	4	93
2	4	62
3	4	95
4	4	66
5	4	52
6	4	65
7	4	97
8	4	66
9	4	91
10	4	113
11	4	109

Adding/Removing Columns

Adding new columns to a `data.frame`: base R

You can add a new column, called `newcol` to `df`, using the `$` operator:

```
df$newcol = df$wt/2.2  
head(df,3)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	1.190909
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	1.306818
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	1.054545

Removing columns to a `data.frame`: base R

You can remove a column by assigning to `NULL`:

```
df$newcol = NULL
```

or selecting only the columns that were not `newcol`:

```
df = df[, colnames(df) != "newcol"]  
head(df, 3)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1

Adding new columns to a `data.frame`: base R

You can also "**column bind**" a `data.frame` with a vector (or series of vectors), using the `cbind` command:

```
cbind(df, newcol = df$wt/2.2)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2

Adding columns to a `data.frame`: `dplyr`

The `mutate` function in `dplyr` allows you to add or replace columns of a `data.frame`:

```
print({df = mutate(df, newcol = wt/2.2)})
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	1.3068182
3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	1.0545455
4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	1.4613636
5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	1.6227273
8	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	1.4500000
9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2	1.4318182
10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
11	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364
12	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	1.8500000
13	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	1.6954545
14	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	1.7181818
15	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	2.3863636
16	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	2.4654545
17	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	2.4295455
18	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1	1.0000000
19	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2	0.7340909

Removing columns to a `data.frame`: `dplyr`

The `NULL` method is still very common.

The `select` function can remove a column with a minus (-), much like removing rows:

```
select(df, -newcol)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
15	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4

Removing columns to a `data.frame`: dplyr

Remove `newcol` and `drat`

```
select(df, -one_of("newcol", "drat"))
```

	mpg	cyl	disp	hp	wt	qsec	vs	am	gear	carb
1	21.0	6	160.0	110	2.620	16.46	0	1	4	4
2	21.0	6	160.0	110	2.875	17.02	0	1	4	4
3	22.8	4	108.0	93	2.320	18.61	1	1	4	1
4	21.4	6	258.0	110	3.215	19.44	1	0	3	1
5	18.7	8	360.0	175	3.440	17.02	0	0	3	2
6	18.1	6	225.0	105	3.460	20.22	1	0	3	1
7	14.3	8	360.0	245	3.570	15.84	0	0	3	4
8	24.4	4	146.7	62	3.190	20.00	1	0	4	2
9	22.8	4	140.8	95	3.150	22.90	1	0	4	2
10	19.2	6	167.6	123	3.440	18.30	1	0	4	4
11	17.8	6	167.6	123	3.440	18.90	1	0	4	4
12	16.4	8	275.8	180	4.070	17.40	0	0	3	3
13	17.3	8	275.8	180	3.730	17.60	0	0	3	3
14	15.2	8	275.8	180	3.780	18.00	0	0	3	3
15	10.4	8	472.0	205	5.250	17.98	0	0	3	4
16	10.4	8	460.0	215	5.424	17.82	0	0	3	4
17	14.7	8	440.0	230	5.345	17.42	0	0	3	4
18	32.4	4	78.7	66	2.200	19.47	1	1	4	1
19	30.4	4	75.7	52	1.615	18.52	1	1	4	2
20	33.9	4	71.1	65	1.835	19.90	1	1	4	1

Ordering columns

Ordering the columns of a `data.frame`: base R

We can use the `colnames` function to get the column names of `df` and then put `newcol` first by subsetting `df` using brackets:

```
cn = colnames(df)
df[, c("newcol", cn[cn != "newcol"])]
```

	newcol	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	1.1909091	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	1.3068182	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	1.0545455	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4	1.4613636	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5	1.5636364	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
6	1.5727273	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7	1.6227273	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8	1.4500000	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9	1.4318182	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10	1.5636364	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11	1.5636364	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12	1.8500000	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13	1.6954545	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14	1.7181818	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
15	2.3863636	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16	2.4654545	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
17	2.4295455	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
18	1.0000000	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1

Ordering the columns of a `data.frame`: `dplyr`

The `select` function can reorder columns. Put `newcol` first, then select the rest of columns:

```
select(df, newcol, everything())
```

	newcol	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	1.1909091	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	1.3068182	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	1.0545455	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4	1.4613636	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5	1.5636364	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
6	1.5727273	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7	1.6227273	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8	1.4500000	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9	1.4318182	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10	1.5636364	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11	1.5636364	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12	1.8500000	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13	1.6954545	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14	1.7181818	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
15	2.3863636	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16	2.4654545	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
17	2.4295455	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
18	1.0000000	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
19	0.7340909	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2

Ordering rows

Ordering the rows of a `data.frame`: base R

We use the `order` function on a vector or set of vectors, in increasing order:

```
df[ order(df$mpg), ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
15	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	2.3863636
16	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	2.4654545
24	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4	1.7454545
7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	1.6227273
17	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	2.4295455
31	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8	1.6227273
14	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	1.7181818
23	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2	1.5613636
22	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2	1.6000000
29	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4	1.4409091
12	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	1.8500000
13	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	1.6954545
11	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364
6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
25	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
30	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	1.3068182

Ordering the rows of a `data.frame`: base R

The `decreasing` argument will order it in decreasing order:

```
df[ order(df$mpg, decreasing = TRUE), ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
20	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1	0.8340909
18	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1	1.0000000
19	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2	0.7340909
28	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2	0.6877273
26	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1	0.8795455
27	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2	0.9727273
8	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	1.4500000
3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	1.0545455
9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2	1.4318182
21	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1	1.1204545
4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	1.4613636
32	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2	1.2636364
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	1.3068182
30	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
25	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
11	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364

Ordering the rows of a `data.frame`: base R

You can pass multiple vectors, and must use the negative (using `-`) to mix decreasing and increasing orderings (sort increasing on `x` and decreasing on `y`):

```
df[ order(df$mpg, -df$hp), ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
16	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	2.4654545
15	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	2.3863636
24	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4	1.7454545
7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	1.6227273
17	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	2.4295455
31	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8	1.6227273
14	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	1.7181818
23	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2	1.5613636
22	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2	1.6000000
29	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4	1.4409091
12	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	1.8500000
13	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	1.6954545
11	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364
6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
25	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
30	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091

Ordering the rows of a `data.frame`: `dplyr`

The `arrange` function can reorder rows By default, `arrange` orders in ascending order:

```
arrange(df, mpg)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
1	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	2.3863636
2	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	2.4654545
3	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4	1.7454545
4	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	1.6227273
5	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	2.4295455
6	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8	1.6227273
7	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	1.7181818
8	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2	1.5613636
9	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2	1.6000000
10	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4	1.4409091
11	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	1.8500000
12	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	1.6954545
13	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364
14	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
15	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
16	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
17	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
18	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
19	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091

Ordering the rows of a `data.frame`: `dplyr`

Use the `desc` to arrange the rows in descending order:

```
arrange(df, desc(mpg))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
1	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1	0.8340909
2	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1	1.0000000
3	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2	0.7340909
4	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2	0.6877273
5	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1	0.8795455
6	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2	0.9727273
7	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	1.4500000
8	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	1.0545455
9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2	1.4318182
10	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1	1.1204545
11	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	1.4613636
12	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2	1.2636364
13	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091
14	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	1.3068182
15	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
16	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
17	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
18	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
19	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
20	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364

Ordering the rows of a `data.frame`: `dplyr`

It is a bit more straightforward to mix increasing and decreasing orderings:

```
arrange(df, mpg, desc(hp))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	newcol
1	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	2.4654545
2	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	2.3863636
3	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4	1.7454545
4	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	1.6227273
5	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	2.4295455
6	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8	1.6227273
7	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	1.7181818
8	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2	1.5613636
9	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2	1.6000000
10	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4	1.4409091
11	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	1.8500000
12	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	1.6954545
13	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	1.5636364
14	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	1.5727273
15	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.5636364
16	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	1.7477273
17	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	1.5636364
18	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	1.2590909
19	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	1.1909091
20	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	1.3068182

Transmutation

The `transmute` function in `dplyr` combines both the `mutate` and `select` functions. One can create new columns and keep the only the columns wanted:

```
transmute(df, newcol2 = wt/2.2, mpg, hp)
```

	newcol2	mpg	hp
1	1.1909091	21.0	110
2	1.3068182	21.0	110
3	1.0545455	22.8	93
4	1.4613636	21.4	110
5	1.5636364	18.7	175
6	1.5727273	18.1	105
7	1.6227273	14.3	245
8	1.4500000	24.4	62
9	1.4318182	22.8	95
10	1.5636364	19.2	123
11	1.5636364	17.8	123
12	1.8500000	16.4	180
13	1.6954545	17.3	180
14	1.7181818	15.2	180
15	2.3863636	10.4	205
16	2.4654545	10.4	215
17	2.4295455	14.7	230
18	1.0000000	32.4	66
19	0.7340909	30.4	52